

A Multi-Agent System for for Impact Story Generation

Jonathan van der Maas, Sem de Lange, Abhinav Agnihotry, Ilyâs el Haroui, Youssef Baccouche

Motivation

Writing Impact Stories now

- ✗ Few researchers actively promote their own work.
- ✗ Time-consuming.
- ✗ Expensive to hire writers.
- ✗ A lot of impact research goes unnoticed and underutilized.

Proposed Solution

A Multi-Agent System

- ✓ Specialized task execution [1].
- ✓ Scalable knowledge processing.
- ✓ Flexible system adaptation [2].
- ✓ Collaboration ensures quality.
- ✓ Integration LLMs in MAS for text writing remains underexplored.

Data

Two Formats

- **Dataset:** 11 UvA stories & corresponding LinkedIn posts.
- **UvA website:** marketing style written post of multiple paragraphs with headers.
- **LinkedIn post:** a relative short story highlighting key findings.
- **ASE & ABS** impact research.

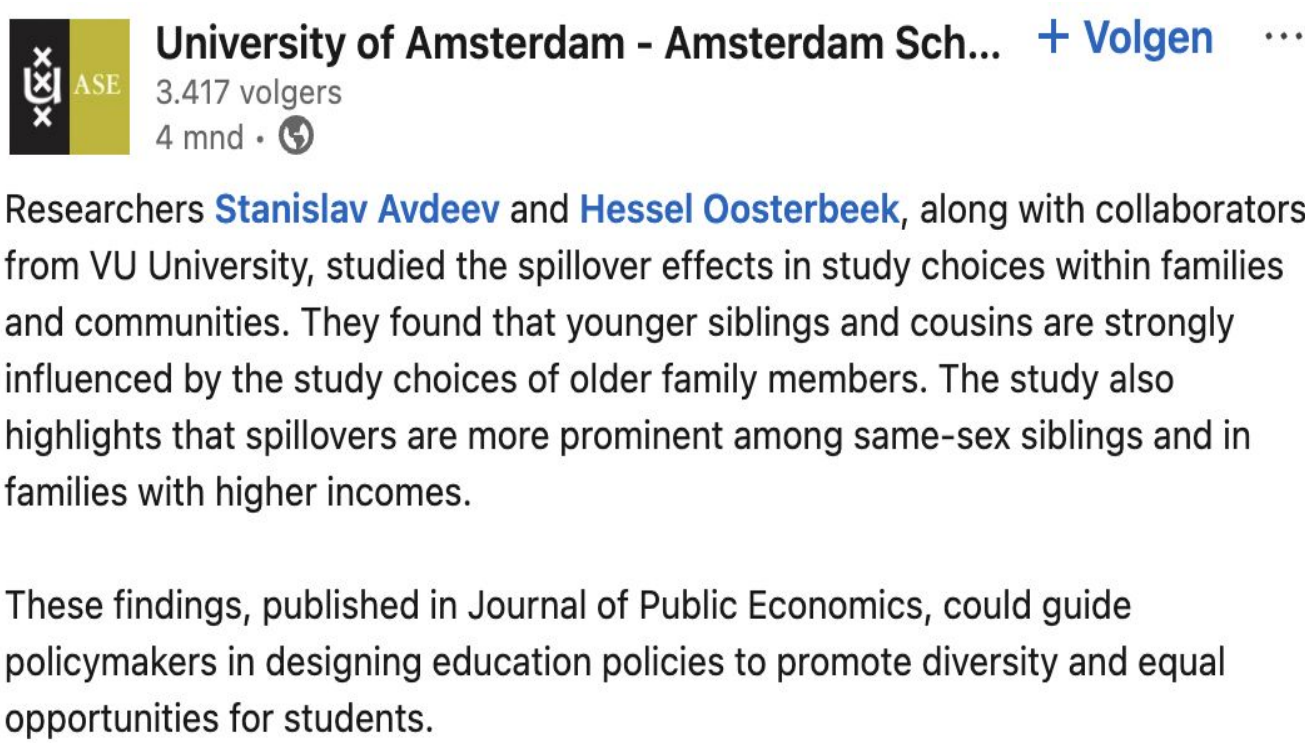


Figure 1. Example LinkedIn post.

Experiments

- Impact stories are generated with MAS and a single LLM to compare performance.
- MAS-generated stories are compared to UvA and LinkedIn references for alignment.
- MAS is tested on research papers of varying lengths and complexities.
- ROUGE, BLEU, compression ratio, and readability scores assess MAS output.

Validation Methods

Quantitative

- ROUGE score
- BLEU score
- Compression Ratio
- FRE

Qualitative

- Stakeholder feedback

Usability

- Flexibility - using edge cases to test.
- Scalability - papers of different sizes and complexities.

The Multi-Agent System

Backend & Preprocessing

Core team outline

- 🤖 **Group chat manager**
 - Organize and moderate team discussions.
- 🤖 **Outline creator**
 - Create structured outlines for impact stories.
- 🤖 **Text writer**
 - Write draft versions of the impact story.
- 🤖 **Text reviewer**
 - Review draft impact stories and send feedback for improvement.
 - Allows for **human feedback**.

Preprocessing

- Data cleaning**
 - Handle non-text elements.
 - Standardize formatting.
- Section segmentation**
 - Divide the research paper into smaller sections.
 - Identify key sections and remove irrelevant sections (e.g., references, acknowledgments).
- Section summarization (BART)**
 - Generate concise summaries of selected sections.
 - Extractive/Abstractive summarization.

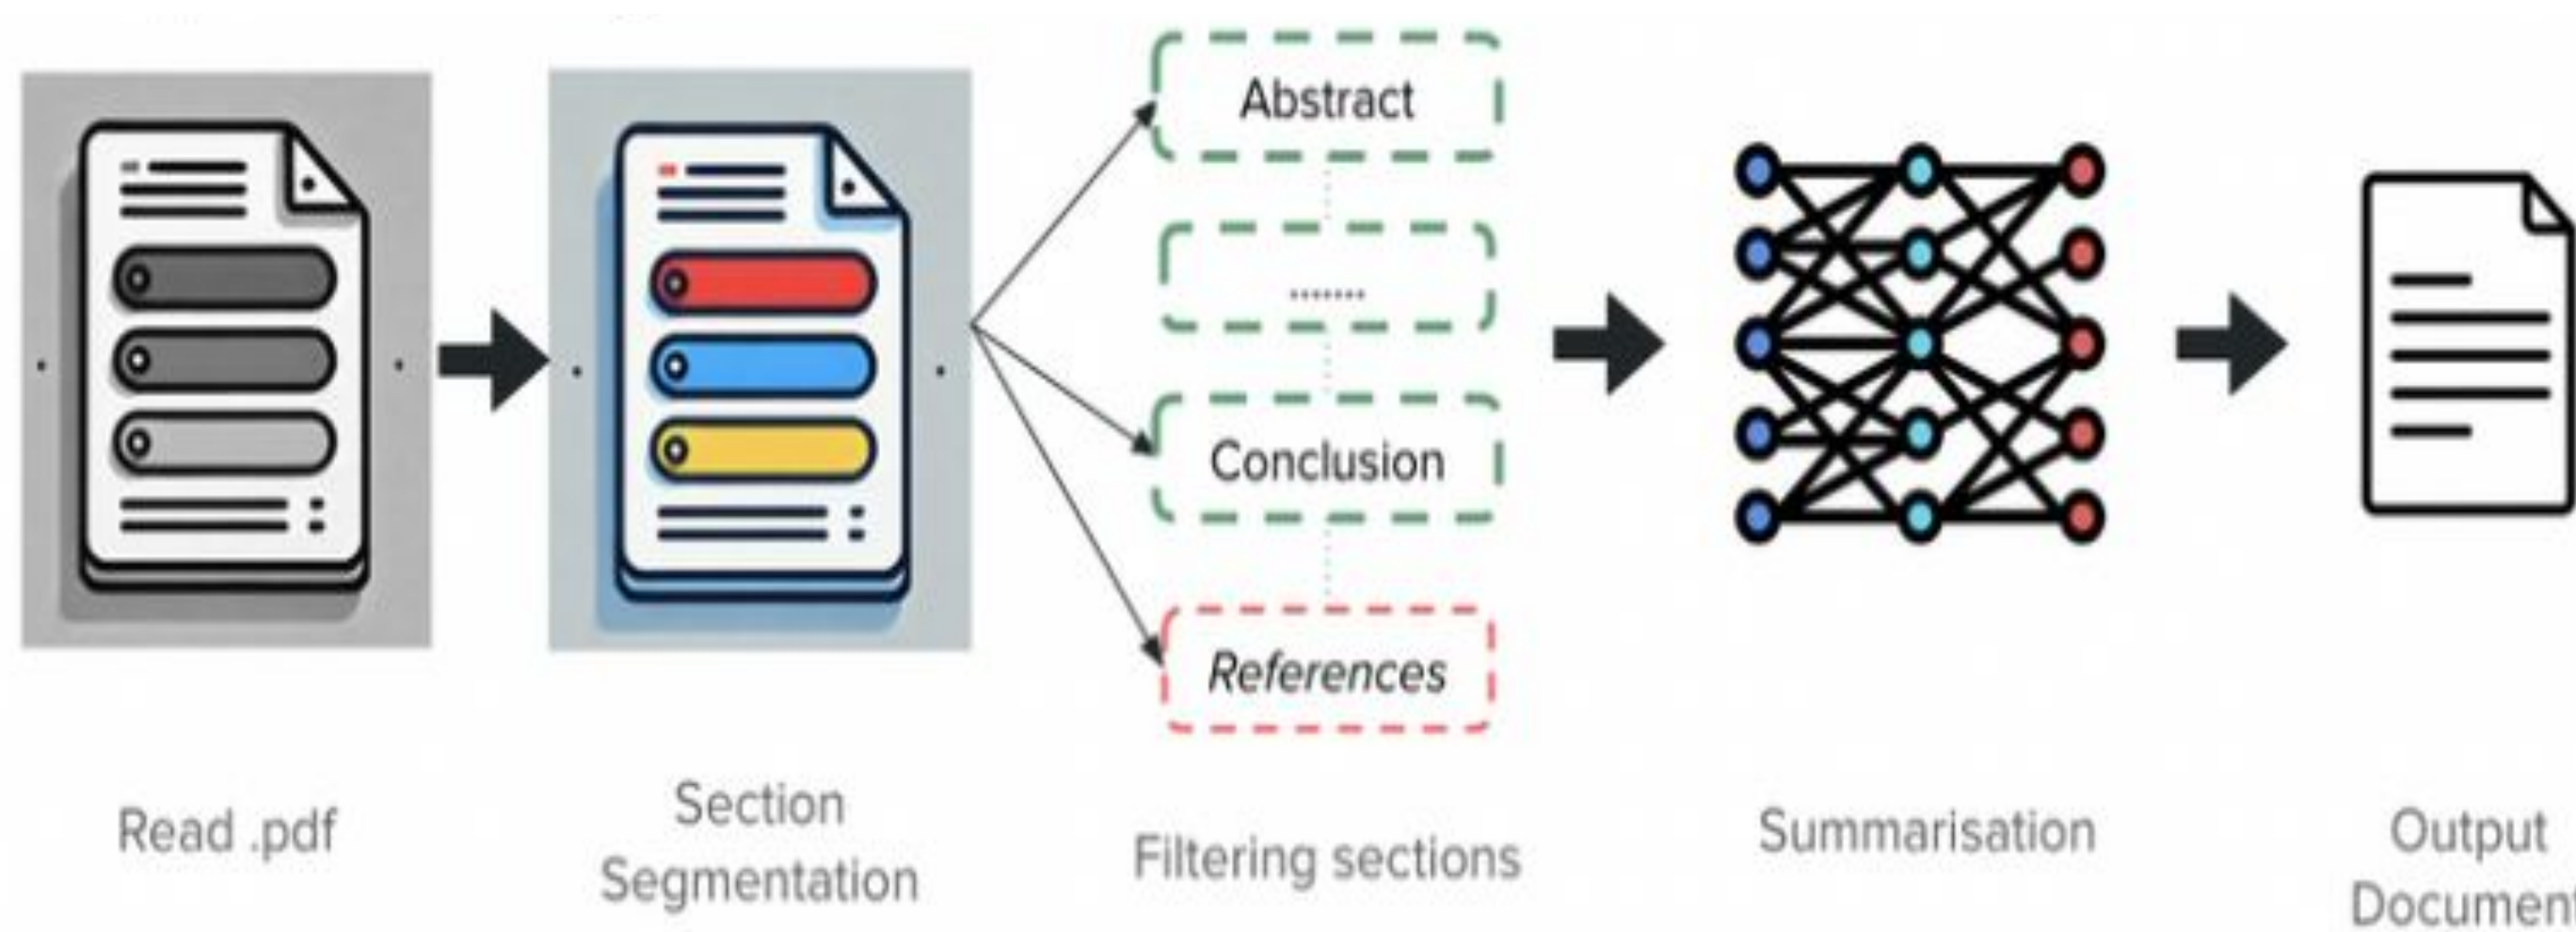


Figure 2. Preprocessing pipeline.

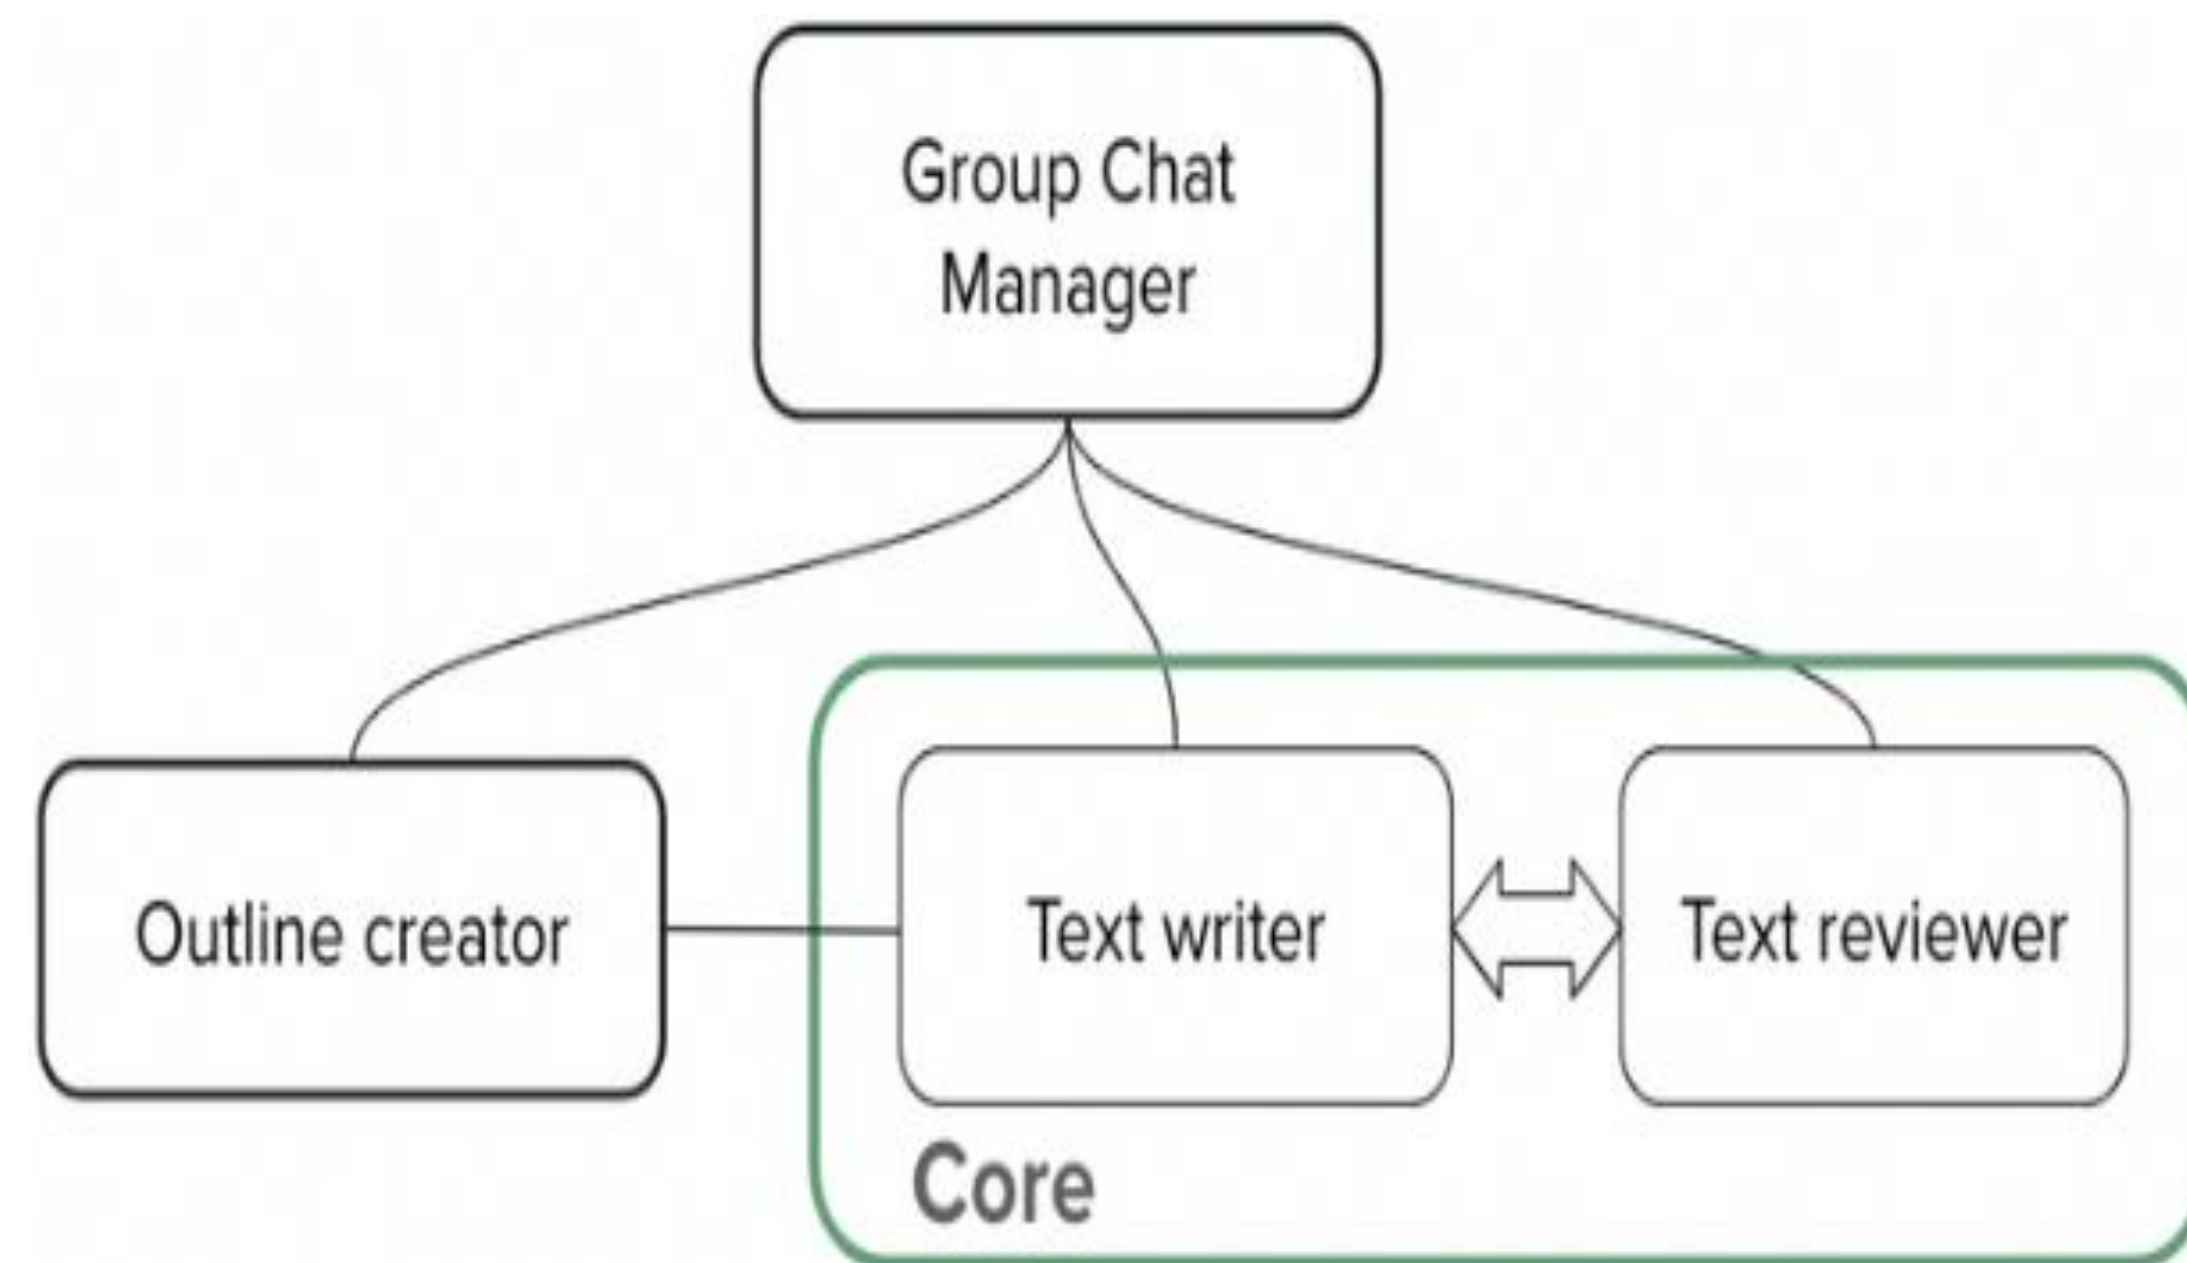


Figure 3. Multi-agent system setup.

Frontend

1. Upload a PDF of the research paper.
2. Select a template (UvA/LinkedIn).
3. Backend preprocesses the research paper and start the multi-agent system workflow.
4. Agents generate the impact story, with the Text Reviewer providing feedback unless a human steps in.
5. Save the draft impact story as a PDF file.

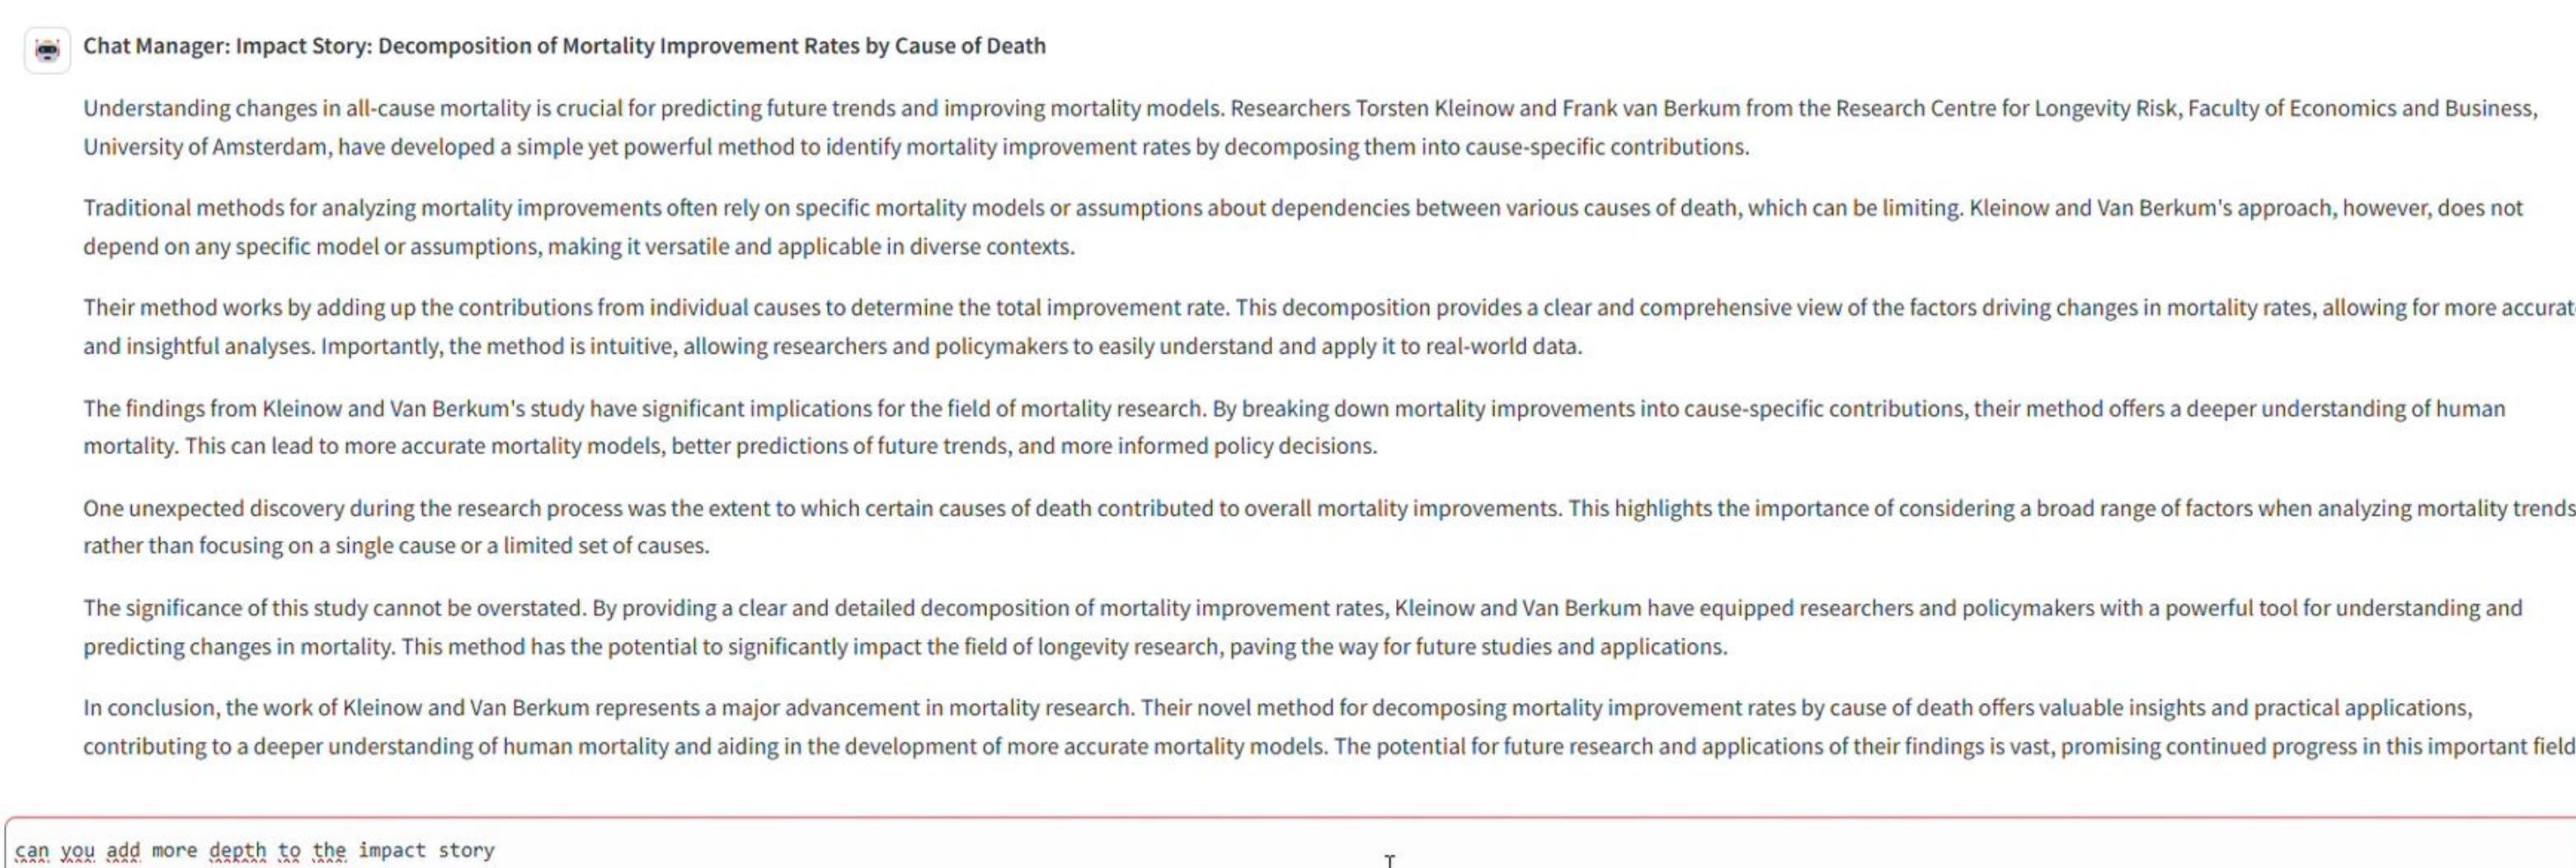


Figure 5. Chat-bot interface with human input.

Results & Conclusion

- The **MAS outperforms a single LLM** agent across **all evaluation metrics**, with better summary alignment and consistency.
- The performance gap is more noticeable in longer templates (UvA).
- The **MAS** provides **more detail and accuracy**, while the single LLM agent struggles with depth.
- Using the MAS offers both a reduction in time spent on generating the impact story as well as the cost associated with it




Figure 6. Comparison MAS generated vs reference UvA stories.

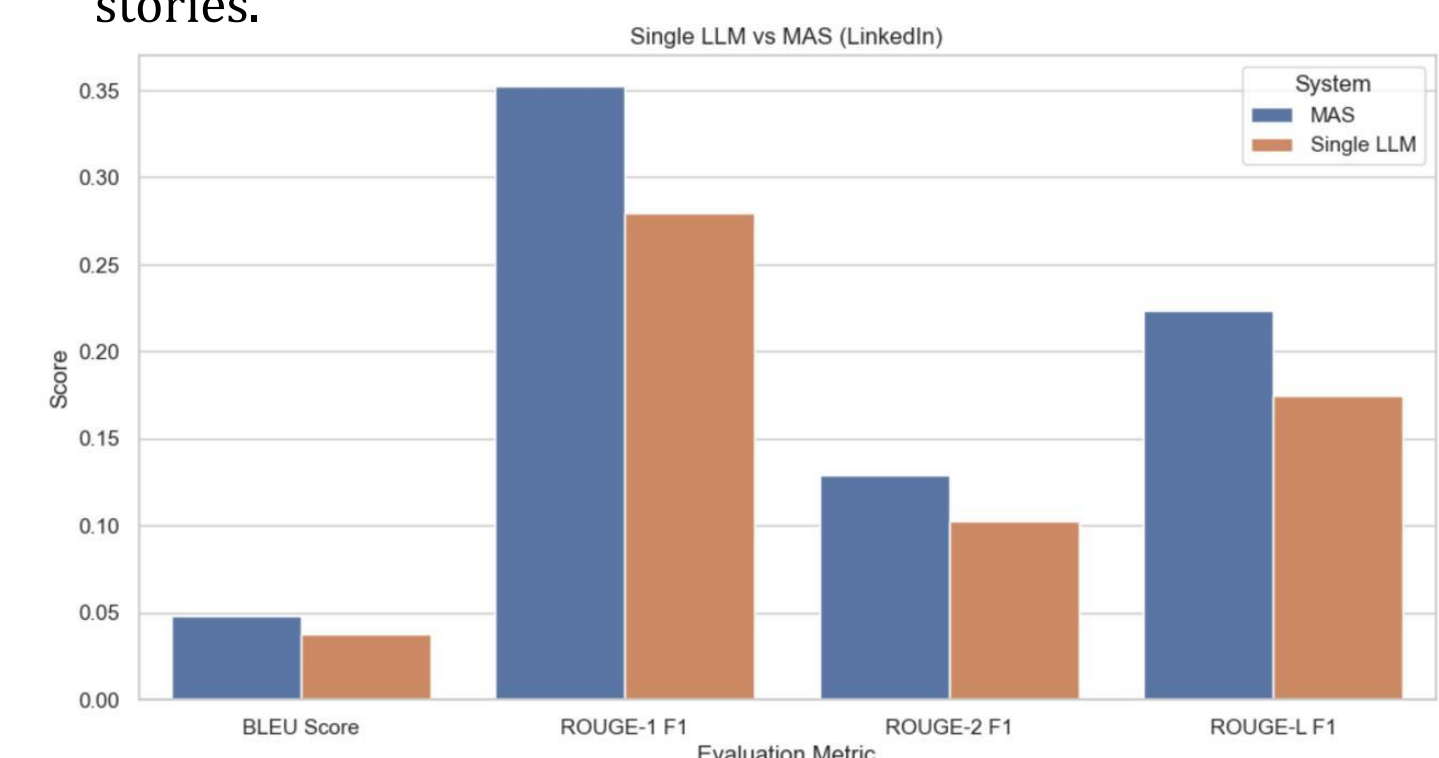


Figure 7. Comparison MAS generated vs reference LinkedIn stories.

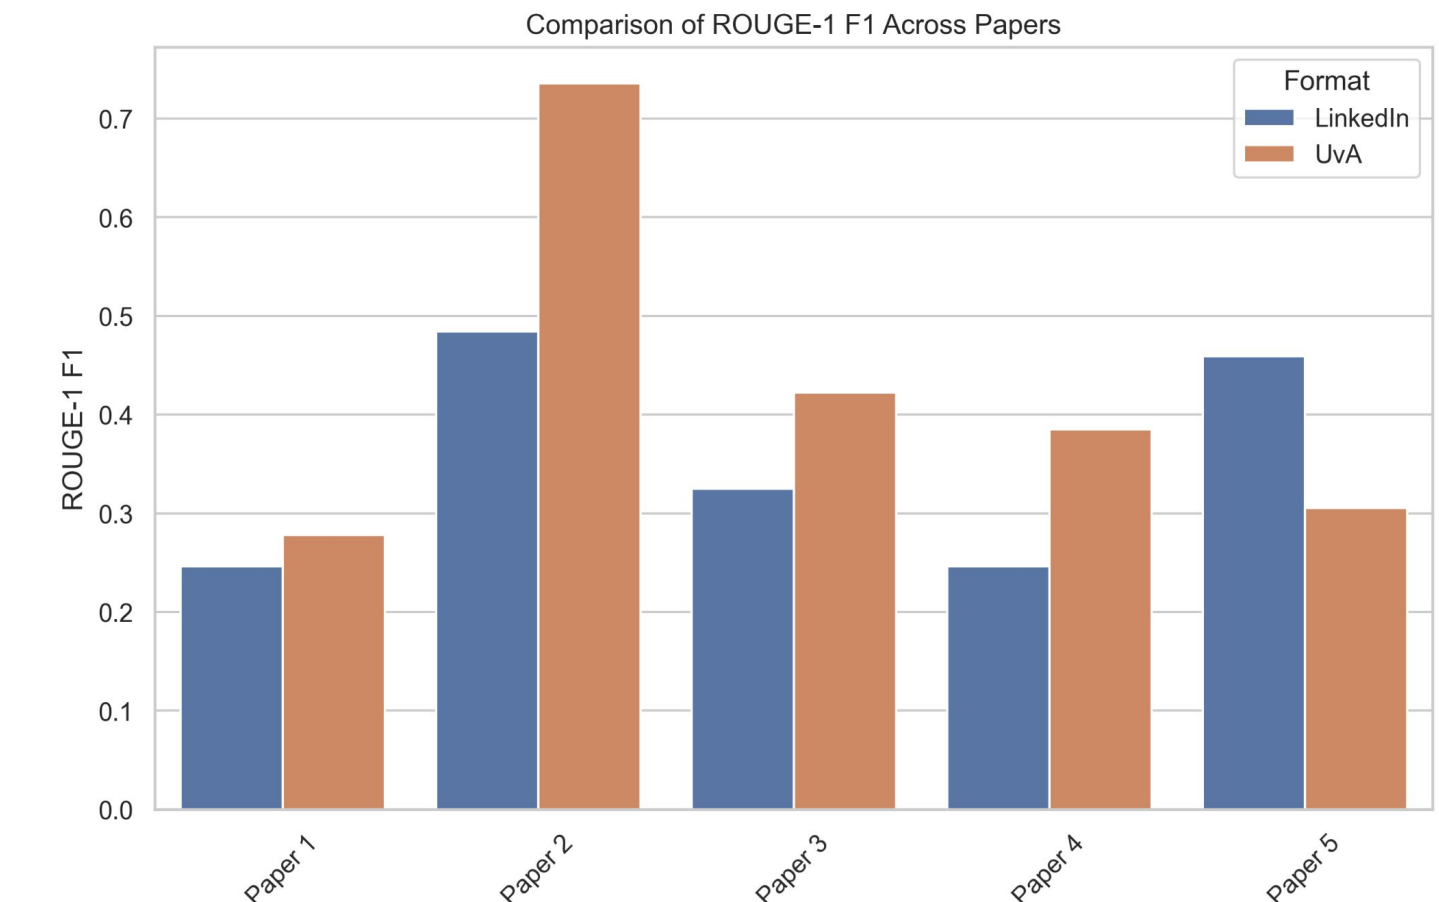


Figure 8. Rouge-1 F1 scores for the different generated stories

Limitations

- API token limits cause in a limit to the conversation length and allowed length of the preprocessed text.
- Difficult to achieve further improvements in generation speed.
- The preprocessing can only handle a single research paper and no other formats such as interviews.

Future work

- Expand format options beyond the two current templates.
- Allow for different types of input documents apart from research papers.
- Enable impact story generation from multiple research papers at once.

References

1. Y. Shoham and K. Leyton-Brown. 2017. Multiagent Systems: Algorithmic, Game-Theoretic, and Logical Foundations. Cambridge University Press.
2. R. F. Da Silva et al . 2019. A Survey of Multi-Agent Systems for Energy Management. Applied Energy 252 (2019), 113491